

# Business Analytics Using Computational Statistics

No Class

Week 4  
Computational Intervals

Week 5  
Bootstrapping

## Functions and Loops

**Math**

Functions define how certain variable  $y$  are associated with other variables  $x$

$$y = f(x)$$

**Programming**

Input: `first_ten <- 1:10`

Parameters: `numbers`

Logic: `standardize <- function(numbers) {  
 numbers <- (numbers - mean(numbers)) / sd(numbers)  
 return(numbers)  
}`

Output: `# [1] -1.4861811 -1.1568120 -0.8257228 ...`

Return Value: `standardize(first_ten)`

Look carefully... `numbers <- (numbers - mean(numbers)) / sd(numbers)`

input vector number is changed in function! Does standardize change the first\_ten vector?

Note: R can also behave like an object-oriented programming language...

## Confidence Intervals

**Confidence Interval of  $\mu$**

(when population statistics are known - unlikely)

Population statistics:  $X \sim N(\mu, \sigma)$

Sample (est. pop.) statistics:  $x \sim N(\bar{x}, s)$

Standard error:  $\sigma_2 = \sigma / \sqrt{n}$

Standardized z-score of sample mean:  $z = \frac{(\bar{x} - \mu)}{\sigma_2}$

Confidence Interval of  $\mu$ :  $\bar{x} \pm z^* \left( \frac{s}{\sqrt{n}} \right)$

95% Confidence Interval:  $\bar{x} - 1.96 \left( \frac{s}{\sqrt{n}} \right)$  to  $\bar{x} + 1.96 \left( \frac{s}{\sqrt{n}} \right)$

"If we took a large number of samples, about 95% of them should contain the population mean in their 95% confidence interval"

| Confidence Level | z    |
|------------------|------|
| 68.3%            | 1.00 |
| 90.0%            | 1.65 |
| 95.0%            | 1.96 |
| 95.0%            | 2.00 |
| 99.0%            | 2.58 |
| 99.7%            | 3.00 |

## Bootstrapping

**Resampling from a Sample**

Population vs. One Sample

Population statistics:  $\mu = 242.4863$

Sample statistics:  $\bar{x} = 240.4981$

Visualizing resampled samples

Plot density of population vs. bootstrapped samples

Plot density of sample vs. bootstrapped samples

# R Packages

*Collections of functions for you to install, load, and use freely*

## Install from CRAN (Comprehensive R Archive Network)

*Mature and well tested packages are usually available on a central repository – CRAN*

```
install.packages("ggplot2")  
  
library("ggplot2")  
# Ready to use all ggplot2 functions
```

## Install from Github

*New or experiental packages are usually available on on Github – where developers self-publish their own code*

```
install.packages("remotes")  
remotes::install_github("soumyaray/compstatslib")  
  
library("compstatslib")  
# Ready to use all compstatslib functions
```

```
machine_precision()  
[1] 2.220446e-16
```

*machine\_precision() reports the smallest possible number on your machine such that  $1 + x \neq 1$*



Soumya Ray



Daniele Melotti

National Tsing Hua University

The screenshot shows the GitHub repository page for 'compstatslib' by 'soumyaray'. The browser address bar shows 'https://github.com/soumyaray/compstatslib'. The page title is 'README.md'. The repository name 'compstatslib' is prominently displayed. Below it, a description reads: 'This R Package is a collection of interactive tools and helper functions that help teachers and learners learn concepts in computational statistics. Useful for homework, in-class demonstrations, or self-learning.'

<https://github.com/soumyaray/compstatslib>

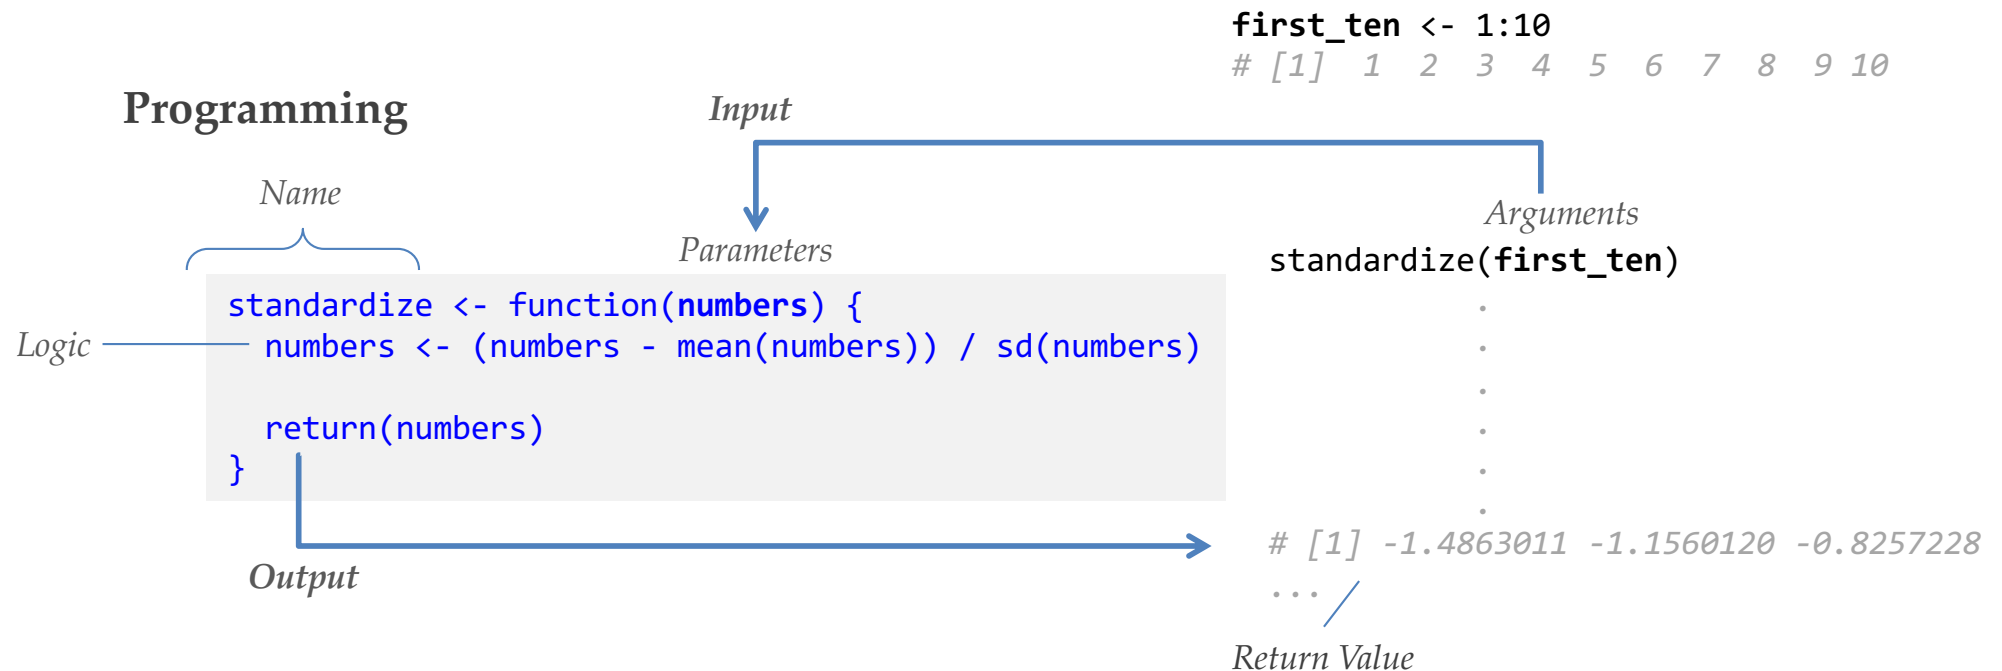
# Functions

## Math

Functions define how certain variable  $y$  are associated with with other variables  $x$

$$y = f(x)$$

## Programming



Look carefully...

```
numbers <- (numbers - mean(numbers)) / sd(numbers)
```

Is the input vector *number*  
is changed in function?




Will `standardize()` change the `first_ten` vector?

# Changing the value of variables

```
standardize <- function(numbers) {  
  numbers <- (numbers - mean(numbers)) / sd(numbers)  
  
  return(numbers)  
}
```

*Look carefully...*

```
numbers <- (numbers - mean(numbers)) / sd(numbers)
```

 *Is the input vector number  
is changed in function?*

 *Will standardize() change the first\_ten vector?*

## Can a function change an input?

```
first_ten <- 1:10
```

```
# [1] 1 2 3 4 5 6 7 8 9 10
```


```
standardized(first_ten)
```

```
# [1] -1.4863011 -1.1560120 -0.8257228 ...
```

```
first_ten
```

```
# [1] 1 2 3 4 5 6 7 8 9 10
```

*It seems the first\_ten vector was **not** changed*

 *What happened when numbers was changed inside the function?*

## Let's see what happens when the value of variables are changed in R:


```
x <- c(1, 2, 3, 4)
```

```
tracemem(x)
[1] "<0x10da9a988>"
```

```
x[2] <- 9
tracemem[0x10da9a988 -> 0x10ddd8a48]:
```

```
x
[1] 1 9 3 4
```

*tracemem() shows us the **address** where a variable's data is stored in computer memory*



## Let's see what happens to variables changed inside functions:


```
standardized_mem <- function(numbers) {
  cat(paste("BEFORE change: numbers is at: ", tracemem(numbers), "\n"))
  numbers <- (numbers - mean(numbers)) / sd(numbers)
  cat(paste("AFTER change: numbers is at: ", tracemem(numbers), "\n"))
  return(numbers)
}
```

```
tracemem(first_ten)
# [1] "<0x7f88405d7e00>"
```

```
stdized <- standardized_mem(first_ten)
# BEFORE change: numbers is at: <0x7f88405d7e00>
# AFTER change: numbers is at: <0x7f8832bb02f8>
```

```
tracemem(stdized)
# [1] "<0x7f8843b6d318>"
```

*R passes objects to functions by reference*



*R does  
"copy on change"  
(makes a new copy of  
data if it is changed)*

# For Loops in R

*A procedural way to repeat operations in code*

*Telling the computer what to do, rather than what you want*

**Problem:** Given a data vector,

Create a vector of boolean (true/false) values to indicate if each data element is positive or not

```
norm_data <- rnorm(500000)
# [1] 0.01874617 -0.18425254 -1.37133055 -0.59916772 0.29454513 0.38979430...

# [1] TRUE FALSE FALSE FALSE TRUE TRUE...
```

**One Solution:** Create a loop that checks each element in the vector of data and makes a new vector of (-1) or (+1) values

*Setup shadow variable*

```
results <- c()
```

*Perform logic*

```
for (i in 1:length(norm_data)) {
  if (norm_data[i] < 0) {
    results[i] <- FALSE
  } else {
    results[i] <- TRUE
  }
}
```

*Build **results** value to return*

*Define **index** range to loop over*

- "i" takes next value in sequence in each iteration

```
# [1] TRUE FALSE FALSE FALSE TRUE TRUE...
```

# The Problem with For-loops

for-loop with *index*

```
results <- c()
for (i in 1:length(norm_data)) {
  if (norm_data[i] < 0) {
    results[i] <- FALSE
  } else {
    results[i] <- TRUE
  }
}

# [1] TRUE FALSE FALSE FALSE TRUE TRUE...
```

**Shadow variables** *make code harder to understand*

**Hard to know the intention of the loop**

*Made for computers to understand*

*Hard for people to understand*

**Easy to write slow, unoptimal code**

*Many optimization tricks must be learned*

*Harder to perform in parallel*

for-loop without *index*

```
results <- c()
for (num in norm_data) {
  if (num < 0) {
    results <- c(results, FALSE)
  } else {
    results <- c(results, TRUE)
  }
}
```



*Which version of the for-loop is easier to understand?  
(try reading each!)*

*Which version of the for loop will be faster? Why?  
(try running each!)*

## Advantages of for-loops?

- More *familiar* to programmers from older languages
- Helpful if *index* of operations is important

# Functional Iteration

*Stating your purpose (function) rather than procedure*

## Apply family (apply/sapply/lapply)

### 1. Define a iteration logic in a **function**

```
is_positive <- function(num) {  
  if (num > 0) {  
    return(TRUE)  
  } else {  
    return(FALSE)  
  }  
}  
  
is_positive(5)  
[1] TRUE
```

#### **Intention is clearer than for-loops**

*We wish to apply the logic of `positive()` function to every element of `norm_data`*

### 2. **Apply** the function over the data

```
results <- sapply(norm_data, is_positive)  
# [1] TRUE FALSE FALSE FALSE TRUE TRUE
```

|        |   |       |
|--------|---|-------|
| 0.018  | → | TRUE  |
| -0.184 | → | FALSE |
| -1.371 | → | FALSE |
| -0.599 | → | FALSE |
| 0.294  | → | TRUE  |

#### **Doesn't have to be optimized like a for-loop**

*R can optimize your functional iterations  
Can be parallelized (see `parallel` package)  
But still not always the fastest or clearest way...*

#### **R treats functions as first-class objects:**

*functions can be stored like data  
functions can be passed as arguments  
functions can even return functions*

*E.g., our function is passed to `sapply()` as an argument*

`sapply(norm_data, is_positive)`



# Vectorized Iteration

*R often already loops through data that is in vector form!*

## Most readable

*Most **operators** are already vectorized*

vector output      vector input      greater-than operator

↓                    ↓                    ↙

```
results <- norm_data > 0
# [1] TRUE FALSE FALSE FALSE TRUE TRUE...
```

*Some **functions** are vectorized as well*

```
ifelse(norm_data > 0, "positive", "negative")
# [1] "positive" "negative" "negative" "negative"
```

## Matches mathematical representation

*Vectorized code is most most likely  
to resemble your math/statistics*

$$x - \bar{x}$$

`x - mean(x)`

$$\frac{\sum (x - \bar{x})^2}{n - 1}$$

`sum((x - mean(x))^2) / (length(x) - 1)`

# Benchmarking Performance

**system.time(...)**

If performance is critical, you can **benchmark** the time performance of your code

*For loops* can be very fast,  
or very slow, depending on  
your experience and the code

```
results <- c();
system.time(
  for (i in 1:length(norm_data)) {
    if (norm_data[i] < 0) { results[i] <- FALSE}
    else { results[i] <- TRUE }
  }
)
# user system elapsed
# 0.159 0.020 0.186
```

*elapsed time is the total time*

*Functional iteration*  
performs somewhere between  
fast and slow for-loops

```
system.time( sapply(norm_data, positive) )
# user system elapsed
# 0.329 0.010 0.340
```

*Vectorized iteration* is  
always fastest, if it is available

```
system.time( norm_data > 0 )
# user system elapsed
# 0.001 0.000 0.001
```

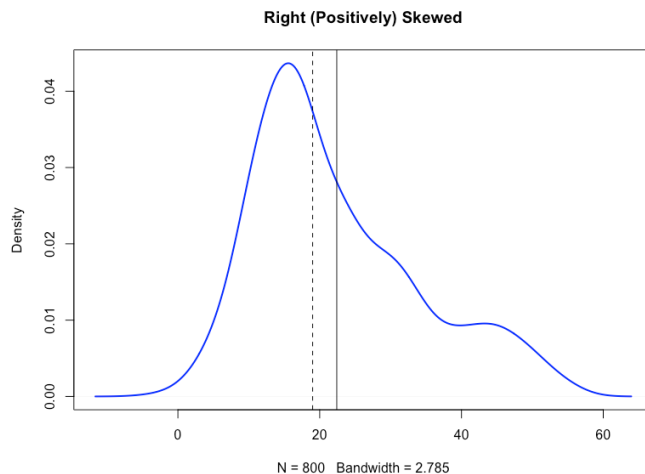
*Making your own function can help make repetitive tasks quicker and easy to change*

```
# Create a visualization function
plot centrality <- function(distr, title) {
  # Plot the full distribution of abc
  plot(density(distr), col="blue", lwd=2, main = title)


  # Add vertical lines showing mean and median
  abline(v=mean(distr))
  abline(v=median(distr), lty="dashed")
}
```

# Mean vs. Median

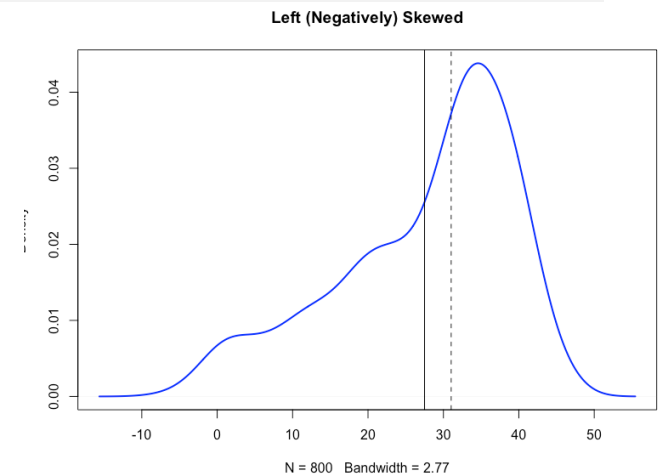
*The mean appears to be more sensitive to outliers*  
*The median (and other quartiles) appears to be more resilient to outliers*



```
d1 <- rnorm(n=500, mean=15, sd=5)
d2 <- rnorm(n=200, mean=30, sd=5)
d3 <- rnorm(n=100, mean=45, sd=5)
d123 <- c(d1, d2, d3)
plot centrality(
  d123,
  title="Right (Positively) Skewed"
)
```



```
d123 <- rnorm(n=800)
plot centrality(d123, title="Normal")
```



```
d1 <- rnorm(n=100, mean=5, sd=5)
d2 <- rnorm(n=200, mean=20, sd=5)
d3 <- rnorm(n=500, mean=35, sd=5)
d123 <- c(d1, d2, d3)
plot centrality(
  d123,
  title="Left (Negatively) Skewed"
)
```




*Why do we always use the **mean** in statistics?*

*Could we do statistics with the **median**?*

# Standard Deviations of a Normal Distribution

```
quartiles_vs_sd <- function(distr) {
  # plot data distribution, mean + standard deviations lines
  plot(density(distr))
  abline(v=mean(distr))
  sd_points <- mean(distr) + (-3:3)*sd(distr)
  abline(v=sd_points, lty='dashed')

  # return the distance of each quartile from the mean
  q = quantile(distr, c(0.25, 0.50, 0.75))
  return((q - mean(distr))/sd(distr))
}
```



## Standard Normal Distribution ( $mean=0$ , $sd=1$ )?

```
data = rnorm(n=10000, mean=0, sd=1)
quartiles_vs_sd(data)
```

| 25%          | 50%          | 75%         |
|--------------|--------------|-------------|
| -0.669003605 | -0.005018409 | 0.678760278 |

Our quartile deviations from the mean are similar if expressed in terms of **standard deviations**

Our quartile deviations from the mean are quite different in terms of their **raw units**


## Non-standard Normal Distribution? ( $sd \neq 1$ )

```
data = rnorm(n=10000, mean=35, sd=3.5)
quartiles_vs_sd(data)
```

| 25%          | 50%         | 75%         |
|--------------|-------------|-------------|
| -0.673143918 | 0.001196753 | 0.672708091 |

### Quartiles and Standard deviation:

- Describes dispersion in normal data, regardless of its center
- Can give us a standard sense of range for normal distributions



# The Empirical Rule


For any symmetrical, bell-shaped distribution

*68% of cases are within 1 sd of the mean*

```
mean(data) + c(-1, 1)*sd(data)
[1] 31.49610 38.50571
```


*95% of cases are within 1.96 sd of the mean*

```
mean(data) + c(-1.96, 1.96)*sd(data)
[1] 28.13150 41.87031
```



*Does not apply for other distributions*

```
> quartiles_vs_sd(d123)
      25%      50%      75%
-0.7513032 -0.3136625  0.6942677
```



*Standard deviation does not describe familiar points in non-normal data*

*Are **quartiles** a better description of dispersion in non-normal data?*

# Interquartile Range Revisited



Should we prefer these metrics to make plots, or should we rely on our own expertise?

**Sturges' Rule:** Log length for number of bins


```
sturges_formula <- function(distr, title) {  
  k = ceiling(log2(length(distr))) + 1  
  h = (max(distr) - min(distr))/k  
  hist(distr, breaks = k)  
  return(data.frame(k, h))  
}
```

*Log length keeps number of bins (k) stable, which is not good when outliers are present, but might be good for other things?*

normal data

```
sturges_formula(rand_data)
```


| k  | h    |
|----|------|
| 11 | 2.60 |



normal + outliers

```
sturges_formula(out_data)
```

| k  | h    |
|----|------|
| 11 | 4.99 |




**Scott's Rule:** Standard deviation for bin size

```
scotts_rule <- function(distr) {  
  h <- 3.5*sd(distr) / (length(distr)^(1/3))  
  k = ceiling((max(distr) - min(distr))/h)  
  hist(distr, breaks = k)  
  return(data.frame(k, h))  
}
```

*Standard deviation is stabler, but could cause problems if extreme outliers are introduced*


```
scotts_rule(rand_data)
```

| k  | h    |
|----|------|
| 15 | 1.92 |



```
scotts_rule(out_data)
```

| k  | h    |
|----|------|
| 24 | 2.32 |




**Freedman-Diaconis Choice:** IQR for bin size

```
fd_choice <- function(distr, title) {  
  h <- 2*IQR(distr) / (length(distr)^(1/3))  
  k <- ceiling((max(distr) - min(distr))/h)  
  hist(distr, breaks = k, main=title)  
  return(data.frame(k, h))  
}
```

*Using quartiles (IQR) makes bin size (h) insensitive to the presence of outliers*


```
fd_choice(norm_data)
```

| k  | h    |
|----|------|
| 20 | 1.49 |



```
fd_choice(rand_data)
```

| k  | h    |
|----|------|
| 38 | 1.48 |



# HW Peer Review Metric

|        | Completeness |   | Solution          |                           | Extra   |
|--------|--------------|---|-------------------|---------------------------|---|
|        | 5            | Everything attempted                                | <i>and</i>        | Everything correct        | <i>and</i> Novel solution or Beautiful report |
| 100% = | <b>4</b>     | <b>Everything attempted</b>                         | <b><i>and</i></b> | <b>Everything correct</b> | <b><i>and</i> Reasonable appearance</b>       |
|        | 3            | Everything attempted                                | <i>and</i>        | Minor mistakes            | <i>or</i> Difficult to interpret              |
|        | 2            | } <b>Missing major parts<br/>or late submission</b> | <i>or</i>         | <b>Major mistakes</b>     |   |
|        | 1            |   |                   |                           |   |
|        | 0            | Not submitted                                       |                   |                           |   |

# HW Suggestions

## **CREATE well formatted reports**

Briefly summarize the question

Format it to distinguish:

*question / description / code / output / answers*

Show code and relevant text output

*use text, not screenshots*

Show relevant visualizations

*export graphics from Rstudio; not screenshots*

## **CREDIT peers who helped**

Mention their ID at the top of your assignment

Peers who help will get extra-credit at end-of-semester

## **REVIEW your peers fairly**

Give specific comments for your response (0-5)

*Why the homework was deducted any points*

*Why the homework was awarded 5 points*

You will get a 'reviewer' grade at end-of-semester:

*accuracy, helpfulness*

## **If you aren't happy with your peer evaluation**

Politely reply to your peer evaluations with a comment

*we will check comments!*

Contact the TAs or professor if no response



# Resampling from a Population

## The Central Limit Theorem

```
library(compstatslib)
```

```
d3 <- rnorm(n=500000, mean=5, sd=5)
d2 <- rnorm(n=200000, mean=20, sd=5)
d1 <- rnorm(n=100000, mean=35, sd=5)
d123 <- c(d1, d2, d3)
interactive_sampling(d123)
```

*distribution of population :*  
(distribution unknown)

Population Mean:  $\mu_x$

Standard Deviation:  $\sigma_x$

*distribution of a sample:*

Sample Mean:

(weakly approx. to pop. mean)

$$\bar{x} = \frac{\sum x_i}{n} \sim \mu_x$$

Standard Deviation:

$$s = \sqrt{\frac{\sum (x_i - \bar{x})^2}{n - 1}}$$

*distribution of all sample means:*

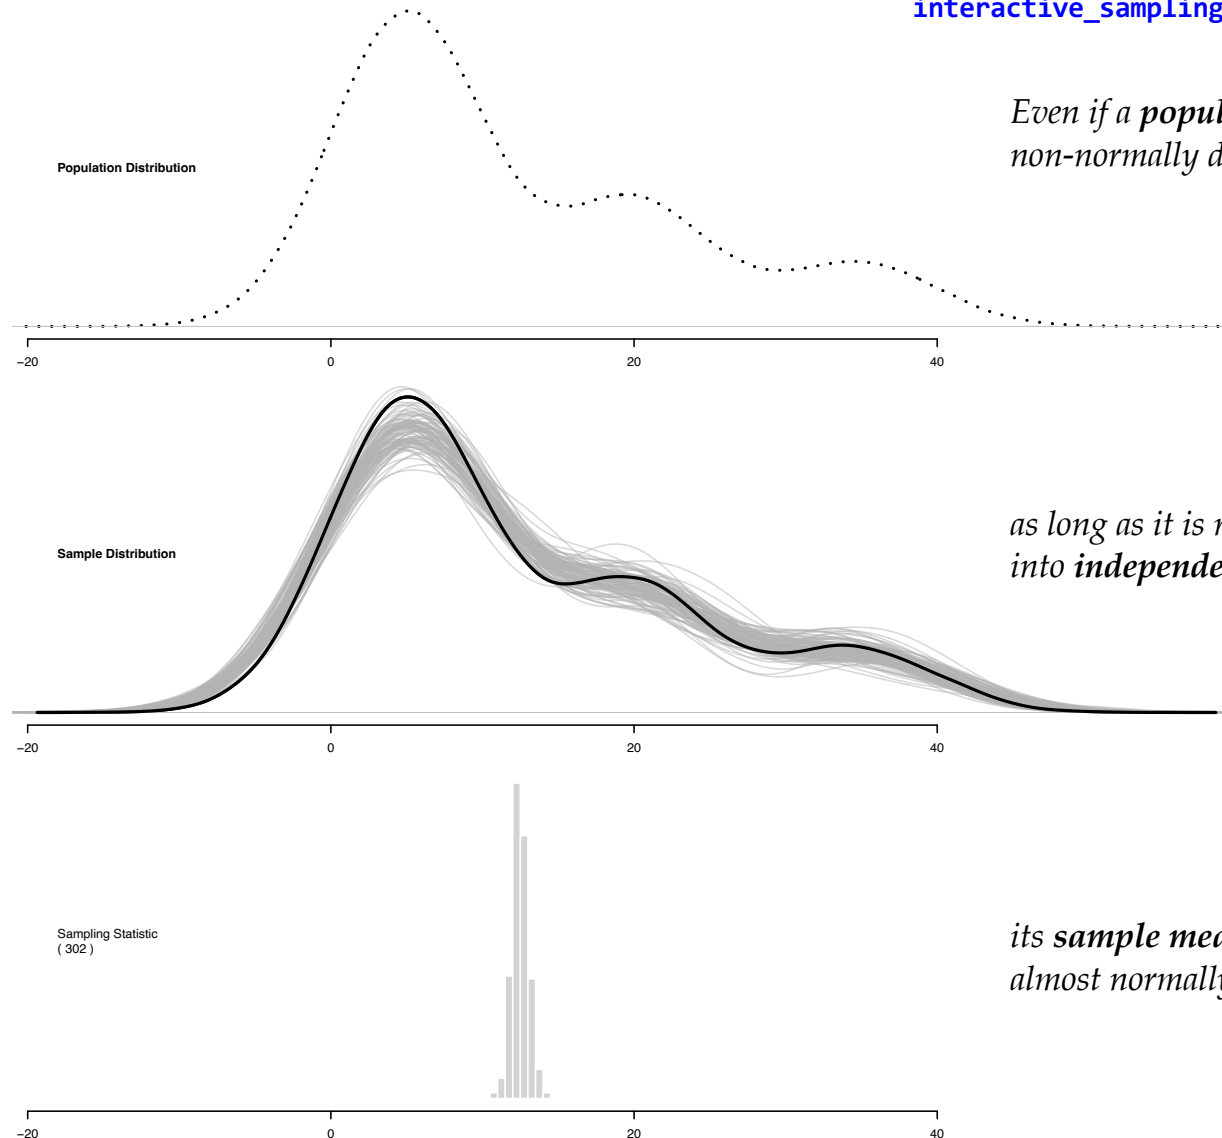
$\bar{x}_1, \bar{x}_2, \bar{x}_3, \dots, \bar{x}_{100}$

Follows a nearly normal distribution

We can approximate its mean and standard deviation with a single sample:

$$\bar{\bar{x}} \sim \bar{x}$$

$$s_x = \frac{s_x}{\sqrt{n}}$$



Even if a **population** is  
non-normally distributed,

as long as it is randomly drawn  
into **independent samples**,

its **sample means** will be  
almost normally distributed.

# Confidence Interval of $\mu$

<https://gist.github.com/soumyaray/285296600b8712b04b52201010b8bd9f>

Population statistics:

*Distribution characteristics unknown*

Sample statistics:

Sample Mean:  
(weakly approx. to pop. mean)

$$\bar{x} = \frac{\sum x_i}{n} \sim \mu_x$$

Standard Deviation:

$$s = \sqrt{\frac{\sum (x_i - \bar{x})^2}{n - 1}}$$

sample size:  $n$

degrees of freedom ( $df$ ) =  $n-1$

**Standard error:**  
(based on one sample)


$$s_{\bar{x}} = \frac{s}{\sqrt{n}}$$

The population mean should be somewhere in the distribution of sampling means

**95% Confidence Interval of the mean:**  $\bar{x} - 1.96 \left( \frac{s}{\sqrt{n}} \right)$  to  $\bar{x} + 1.96 \left( \frac{s}{\sqrt{n}} \right)$

**99% Confidence Interval of the mean:**  $\bar{x} - 2.58 \left( \frac{s}{\sqrt{n}} \right)$  to  $\bar{x} + 2.58 \left( \frac{s}{\sqrt{n}} \right)$

distribution of sample means  $\bar{x}_1, \bar{x}_2, \bar{x}_3, \dots, \bar{x}_{100}$



**Confidence Interval of Population Mean ( $\mu_x$ ):**

$$\bar{x} \pm t \left( \frac{s}{\sqrt{n}} \right)$$


| Confidence Level | t ( $df > 30$ ) |
|------------------|-----------------|
| 90.0%            | 1.65            |
| 95.0%            | 1.96            |
| 99.0%            | 2.58            |

If we take a large number of samples,

~95% of samples should contain the population mean in their 95% confidence interval,

~99% of samples should contain the population mean in their 99% confidence interval

```
library(compstatslib)
plot_sample_ci()
```



# Resampling with Replacement

## Resampling from a Uniform Distribution:

```
seq_sample <- 1:22  
resample <- sample(seq_sample, replace=TRUE)
```

```
seq_sample  
[1] 1 2 3 4 5 6 7 8 9 10 11 12 13 14 15 16 17 18 19 20 21 22  
sort(resample)  
[1] 1 1 3 4 6 8 9 9 13 13 13 14 14 15 15 16 17 20 20 20 20 22
```



*Recall: we learned how to "sample with replacement"  
Is this data simulation helpful?*


## Resampling from a Bell-shaped Distribution

```
rep_sample <- c(1, rep(2,2), rep(3,3), rep(4,4), rep(5,3), rep(6,2), 7)  
resample <- sample(rep_sample, replace=TRUE)
```

```
rep_sample  
[1] 1 2 2 3 3 3 4 4 4 4 5 5 5 6 6 7
```

```
sort(resample)  
[1] 1 1 2 3 3 4 4 4 4 5 5 5 6 6 6 7
```

```
plot(density(rep_sample), lwd=3)  
lines(density(resample), lty="dashed")
```



*If we randomly pick elements from a sample with replacement,  
our new sample will have a **similar distribution** to our original sample*

# Population vs. Sample

## The Unseen Population


*Imagine a population that we cannot measure*

```
a <- rnorm(n=100000, mean=150, sd=15)
b <- rnorm(n=200000, mean=190, sd=25)
c <- rnorm(n=500000, mean=255, sd=25)
d <- rnorm(n=200000, mean=310, sd=20)
```

```
population <- c(a,b,c,d)
pop_mean <- mean(population)
```

```
plot(density(population), col="blue", lty="dashed")
abline(v=pop_mean, lty="dashed")
```

```
pop_mean
# [1] 242.5034
```



## Taking One Sample

*Let's take a sample (sample0) from the population: it's the only thing we can measure*

```
sample_size = 300
sample0 = sample(pop, sample_size)

sample0_mean = mean(sample0)
```

*The sample will not have exactly the same distribution or descriptives as the population*

```
lines(density(sample0), col="blue", lwd=2)
abline(v=sample0_mean, lwd=2)
```

```
sample0_mean
# [1] 239.6902
```



*But resampling loses much of the information about the population*

*Can we say anything about the population?*

# The Bootstrap: Computational Resampling for Inference

We can *resample* from our *original sample* to see where most of the sampling means are.

To start a process with no input or help  
"bootup the computer"  
"pick yourself up by your bootstraps!"

## Bootstrapped resamples

```
resamples <- replicate(3000,  
  sample(sample0, length(sample0), replace=TRUE))
```

`replicate(n, expr):`  
repeat an operation *n* times

```
dim(resamples)  
[1] 300 3000
```


*replicate* returns 300 rows (data in each sample) and 3000 columns (samples)

## Visualizing our resampled samples

```
# Create an empty plotting space with axes  
plot(density(population), lwd=0, ylim=c(0, 0.009))  
  
# A function to plot a single sample's distribution  
plot_resample_density <- function(sample_i) {  
  lines(density(sample_i), col=rgb(0.0, 0.4, 0.0, 0.01))  
  return(mean(sample_i))  
}  
  
# Iteratively plot and get means of all bootstrapped samples  
sample_means <- apply(resamples, 2, FUN=plot_resample_density)  
  
# Plot hidden population and original sample distributions  
lines(density(sample0), lwd=3)  
lines(density(population), lwd=2, lty="dashed")
```

`plot_resample_density()` will plot and compute mean for a single resample

population vs. bootstrapped samples




`rgb(0.0, 0.4, 0.0, 0.01)`  
red green blue alpha  
transparency

## Visualizing Resampled Means

```
# Plot population and original sample densities
plot(density(pop), col="blue", lty="dashed")
lines(density(sample0), col="blue", lwd=2)

# Draw light vertical lines for each sampling mean
abline(v=sample_means, col=rgb(0.7, 0.7, 0.7, 0.01))

# Draw dark lines of population and original sample mean
abline(v=mean(sample_means), lwd=2)
abline(v=pop_mean, lty="dashed")
```



## Distribution of Resampled Means

```
## Distribution of sampling
plot(density(sample_means), lwd=2, xlim=c(0, 400))

## Confidence intervals of the sampling means
quantile(sample_means, probs=c(0.025, 0.975))
# 2.5% 97.5%
# 233.7158 245.9080

quantile(sample_means, probs=c(0.005, 0.995))
# 0.5% 99.5%
# 231.8270 247.8224
```

