

HW16

111078513

2023-06-02

```
# loads some packages we need
library(rpart)
library(rpart.plot)

# Load the data and remove missing values
cars <- read.table("auto-data.txt", header=FALSE, na.strings = "?")
names(cars) <- c("mpg", "cylinders", "displacement", "horsepower", "weight",
               "acceleration", "model_year", "origin", "car_name")
cars$car_name <- NULL
cars <- na.omit(cars)

# IMPORTANT: Shuffle the rows of data in advance for this project!
set.seed(27935752)
cars <- cars[sample(1:nrow(cars)),]

# DV and IV of formulas we are interested in
cars_full <- mpg ~ cylinders + displacement + horsepower + weight +
  acceleration + model_year + factor(origin)
cars_reduced <- mpg ~ weight + acceleration + model_year + factor(origin)
cars_full_poly2 <- mpg ~ poly(cylinders, 2) + poly(displacement, 2) + poly(horsepower, 2) +
  poly(weight, 2) + poly(acceleration, 2) + model_year +
  factor(origin)
cars_reduced_poly2 <- mpg ~ poly(weight, 2) + poly(acceleration, 2) + model_year +
  factor(origin)
cars_reduced_poly6 <- mpg ~ poly(weight, 6) + poly(acceleration, 6) + model_year +
  factor(origin)
```

Question 1) Compute and report the in-sample fitting error (MSE_{in}) of all the models described above. It might be easier to first write a function called `mse_in(...)` that returns the fitting error of a single model.

```
# Run the models.
attach(cars)
lm_full <- lm(cars_full)
lm_reduced <- lm(cars_reduced)
lm_poly2_full <- lm(cars_full_poly2)
lm_poly2_reduced <- lm(cars_reduced_poly2)
lm_poly6_reduced <- lm(cars_reduced_poly6)
rt_full <- rpart(cars_full)
rt_reduced <- rpart(cars_reduced)
```

```

# Make the list we'll need
model_list <- list(lm_full, lm_reduced, lm_poly2_full, lm_poly2_reduced,
                  lm_poly6_reduced, rt_full, rt_reduced)
name_list <- c("lm(cars_full)", "lm(lm_reduced)", "lm(cars_full_poly2)",
              "lm(cars_reduced_poly2)", "lm(cars_reduced_poly6)",
              "rpart(cars_full)", "rpart(cars_reduced)")
formula_list <- c(cars_full, cars_reduced, cars_full_poly2,
                 cars_reduced_poly2, cars_reduced_poly6,
                 cars_full, cars_reduced)
function_list <- c(lm, lm, lm, lm, lm, rpart, rpart)

# Build the function we'll need.
mse_in <- function(model){
  mean(residuals(model)^2)
}

# Get all MSE of models above.
MSE_in_list <- sapply(1:7, \(i){
  model <- model_list[[i]]
  mse_in <- mse_in(model)
})
names(MSE_in_list) <- name_list
MSE_in_list

##           lm(cars_full)           lm(lm_reduced)           lm(cars_full_poly2)
##           10.682122           10.971643           7.919030
## lm(cars_reduced_poly2) lm(cars_reduced_poly6)           rpart(cars_full)
##           8.364546           8.254377           9.155146
##           rpart(cars_reduced)
##           9.501344

```

Question 2) Let's try some simple evaluation of prediction error. Let's work with the `lm_reduced` model and test its predictive performance with split-sample testing:

a. Split the data into 70:30 for training:test.

```

train_set <- cars[1:(nrow(cars)*0.7), ]
test_set <- cars[-(1:(nrow(cars)*0.7)), ]

```

b. Retrain the `lm_reduced` model on just the training dataset (call the new model: `trained_model`); Show the coefficients of the trained model.

```

train_model <- lm(cars_reduced, data = train_set)
summary(train_model)$coefficients

##           Estimate   Std. Error   t value   Pr(>|t|)
## (Intercept) -25.18094941  5.0090298892  -5.0271110 9.119436e-07
## weight      -0.00551164  0.0003394218 -16.2383237 9.347739e-42
## acceleration  0.06855571  0.0855923713   0.8009559 4.238667e-01
## model_year   0.82885622  0.0631211683  13.1311926 9.537389e-31
## factor(origin)2  3.22722498  0.6527806111   4.9438125 1.352589e-06

```

```
## factor(origin)3    2.03952491 0.6289446776    3.2427731 1.333655e-03
```

c. Use the trained_model model to predict the mpg of the test dataset

(i) What is the in-sample mean-square fitting error (MSE_{in}) of the trained model?

```
mean(residuals(train_model)^2)
```

```
## [1] 11.70576
```

(ii) What is the out-of-sample mean-square prediction error (MSE_{out}) of the test dataset?

```
pred <- predict(train_model, test_set)
mean((test_set$mpg - pred)^2)
```

```
## [1] 10.2117
```

d. Show a data frame of the test set's actual mpg values, the predicted mpg values, and the difference of the two (ε_{out} = predictive error); Just show us the first several rows of this dataframe.

```
mpg_actual_pred <- data.frame("pred" = pred, "actual" = test_set$mpg)
head(mpg_actual_pred, 5)
```

```
##      pred actual
## 89  14.04156   14.0
## 215 17.51041   13.0
## 288 19.40509   16.5
## 179 25.07802   23.0
## 398 29.12358   31.0
```

Question 3) Let's use k-fold cross validation (k-fold CV) to see how all these models perform predictively!

a. Write a function that performs k-fold cross-validation (see class notes and ask us online for hints!). Name your function `k_fold_mse(model, dataset, k=10, ...)` – it should return the MSE_{out} of the operation. Your function must accept a model, dataset and number of folds (k) but can also have whatever other parameters you wish.

```
# Calculate prediction error for fold i out of k
# n for nrow(dataset), formula for formula, func for lm() or rpart()
fold_i_pe <- function(i, k, dataset, n, formula, func) {
  folds <- cut(1:n, breaks = k, labels = FALSE)
  test_indices <- which(folds == i)
  test_set <- dataset[test_indices, ]
  train_set <- dataset[-test_indices, ]
  trained_model <- func(formula, data = train_set)
  predictions <- predict(trained_model, test_set)
  test_set$mpg - predictions
}
```

```
# Calculate mse_out across all folds
k_fold_mse <- function(dataset, k=10, formula, func) {
  fold_pred_errors <- sapply(1:k, \i) {
    fold_i_pe(i, k, dataset, nrow(dataset), formula, func)}
  pred_errors <- unlist(fold_pred_errors)
  mean(pred_errors^2)
}
```

(i) Use your `k_fold_mse` function to find and report the 10-fold CV MSE_{out} for all models.

```
MSE_out_list <- sapply(1:7, \i){k_fold_mse(cars, k = 10,
                                          formula = formula_list[[i]],
                                          func = function_list[[i]])}
names(MSE_out_list) <- name_list
MSE_out_list
```

##	lm(cars_full)	lm(lm_reduced)	lm(cars_full_poly2)
##	11.262460	11.415855	8.599373
##	lm(cars_reduced_poly2)	lm(cars_reduced_poly6)	rpart(cars_full)
##	8.818607	9.267369	13.342221
##	rpart(cars_reduced)		
##	13.476272		

(ii) For all the models, which is bigger — the fit error (MSE_{in}) or the prediction error (MSE_{out})?

```
MSE_df <- data.frame(method = name_list, MSE_in = MSE_in_list,
                     MSE_out = MSE_out_list)
MSE_df[, c("MSE_in", "MSE_out")]
```

##		MSE_in	MSE_out
##	lm(cars_full)	10.682122	11.262460
##	lm(lm_reduced)	10.971643	11.415855
##	lm(cars_full_poly2)	7.919030	8.599373
##	lm(cars_reduced_poly2)	8.364546	8.818607
##	lm(cars_reduced_poly6)	8.254377	9.267369
##	rpart(cars_full)	9.155146	13.342221
##	rpart(cars_reduced)	9.501344	13.476272

The MSE_{out} values of all models are higher than the MSE_{in} values.

(iii) Does the 10-fold MSE_{out} of a model remain stable (same value) if you re-estimate it over and over again, or does it vary?

```
tentimes_kfold <- sapply(1:7, \i){
  mean(replicate(10, k_fold_mse(cars[sample(1:nrow(cars))],
                                k = 10, formula = formula_list[[i]],
                                func = function_list[[i]])))
MSE_df$"10 times k-fold MSE_out" = tentimes_kfold
MSE_df[, c(3, 4)]
```

##		MSE_out	10 times k-fold MSE_out
##	lm(cars_full)	11.262460	11.321864
##	lm(lm_reduced)	11.415855	11.403597

```
## lm(cars_full_poly2)      8.599373      8.645329
## lm(cars_reduced_poly2)  8.818607      8.861234
## lm(cars_reduced_poly6)  9.267369      9.285558
## rpart(cars_full)        13.342221     12.816573
## rpart(cars_reduced)     13.476272     12.716515
```

b. Make sure your `k_fold_mse()` function can accept as many folds as there are rows (i.e., `k=392`).

(i) How many rows are in the training dataset and test dataset of each iteration of k-fold CV when `k=392`?

```
n = nrow(cars)
k = 392
folds <- cut(1:n, breaks = k, labels = FALSE)
test_indices <- which(folds == 1)
train_set <- cars[-test_indices, ]
print(paste("There are ", nrow(train_set), " in train_set", sep = ""))

## [1] "There are 391 in train_set"

test_set <- cars[test_indices, ]
print(paste("There are ", nrow(test_set), " in test_set", sep = ""))

## [1] "There are 1 in test_set"
```

(ii) Report the k-fold CV MSE_{out} for all models using `k=392`.

```
MSE_out_list_392 <- sapply(1:7, \(i)
  {k_fold_mse(cars, k = 392,
               formula = formula_list[[i]],
               func = function_list[[i]])})
MSE_df$"MSE_out_392" <- MSE_out_list_392
names(MSE_out_list_392) <- name_list
MSE_out_list_392

##          lm(cars_full)      lm(lm_reduced)      lm(cars_full_poly2)
##          11.293439          11.380040          8.610385
## lm(cars_reduced_poly2) lm(cars_reduced_poly6)      rpart(cars_full)
##          8.787013          9.177932          12.769791
##      rpart(cars_reduced)
##          13.145150
```

(iii) When `k=392`, does the MSE_{out} of a model remain stable (same value) if you re-estimate it over and over again, or does it vary? (show a few repetitions for any model and decide!)

```
tentimes_kfold_392 <- sapply(1:7, \(i){
  mean(replicate(10, k_fold_mse(cars[sample(1:nrow(cars)), ],
                                k = 392, formula = formula_list[[i]],
                                func = function_list[[i]])))})
MSE_df$"10 times k-fold MSE_out_392" <- tentimes_kfold_392
names(tentimes_kfold_392) <- name_list
tentimes_kfold_392

##          lm(cars_full)      lm(lm_reduced)      lm(cars_full_poly2)
```

```
##           11.293439           11.380040           8.610385
## lm(cars_reduced_poly2) lm(cars_reduced_poly6)      rpart(cars_full)
##           8.787013           9.177932           12.769791
##      rpart(cars_reduced)
##           13.145150
```

(iv) Looking at the fit error (MSE_{in}) and prediction error (MSE_{out} ; $k=392$) of the full models versus their reduced counterparts (with the same training technique), does multicollinearity present in the full models seem to hurt their fit error and/or prediction error?

```
MSE_df[c(1, 2, 6, 7), c("MSE_in", "10 times k-fold MSE_out_392")]
```

```
##           MSE_in 10 times k-fold MSE_out_392
## lm(cars_full)      10.682122           11.29344
## lm(lm_reduced)     10.971643           11.38004
## rpart(cars_full)    9.155146           12.76979
## rpart(cars_reduced) 9.501344           13.14515
```

Although the differences may seem small, both in the `lm()` and `rpart()` cases, the reduce model's MSE_{in} and MSE_{out} are slightly higher than those of the full model.

(v) Look at the fit error and prediction error ($k=392$) of the reduced quadratic versus 6th order polynomial regressions — did adding more higher-order terms hurt the fit and/or predictions?

```
MSE_df[c(4, 5), c("MSE_in", "10 times k-fold MSE_out_392")]
```

```
##           MSE_in 10 times k-fold MSE_out_392
## lm(cars_reduced_poly2) 8.364546           8.787013
## lm(cars_reduced_poly6) 8.254377           9.177932
```

In the case of the sixth-degree polynomial model, the MSE_{in} is slightly lower compared to the quadratic polynomial model. However, in terms of MSE_{out} , the quadratic polynomial model has a lower value.